

# CV2\_HW4\_dh3027

April 18, 2022

## 0.1 GANs : Generative Adversarial Networks

Image from [here](#)

A generative adversarial network (GAN) is a generative model composed of two neural networks: a generator and a discriminator. These two networks are trained in unsupervised way via competition. The generator creates "realistic" fake images from random noise to fool the discriminator, while the discriminator evaluates the given image for authenticity. The loss function that the generator wants to minimize and the discriminator to maximize is as follows:

$$\min_G \max_D L(D, G) = \mathbb{E}_x p_{data}(x) [\log D(x)] + \mathbb{E}_z p_z(z) [\log(1 - D(G(z)))]$$

Here,  $G$  and  $D$  are the generator and the discriminator. The first and second term of the loss represent the correct prediction of the discriminator on the real images and on the fake images respectively.

## 0.2 DCGAN

- You will implement deep convolutional GAN model on the MNIST dataset with Pytorch. The input image size is 28 x 28.
- The details of the generator of DCGAN is described below.
- You will start with batch size of 128, input noise of 100 dimension and Adam optimizer with learning rate of 2e-4. You may vary these hyperparameters for better performance.

## 0.3 Architectures

Generator:

The goal for the generator is to use layers such as convolution, maybe also upsampling layer/transposedConvolution to produce image from the given input noise vector. As this is DCGAN (deep convolutional GAN), we expect you to use convolution in the generator. You will get full credit if you can produce [batchsize, 1, 28, 28] vector (image) from the given [batchsize, 100, 1, 1] vector (noise).

Linear Layers that you may use:

- torch.nn.Conv2d
- torch.nn.UpsamplingBilinear2d
- torch.nn.ConvTranspose2d

Non-linear layer:

- `torch.nn.LeakyReLU` with `slope=0.2` between all linear layers.
- `torch.nn.Tanh` for the last layer's activation. Can you explain why do we need this in the code comment?

You may use `view` to change the vector size: <https://pytorch.org/docs/stable/generated/torch.Tensor.view.html>

We recommend to use 2 Conv/TransposedConv layers. When you are increasing the feature map size, considering upsample the feature by a factor of 2 each time. If you have width of 7 in one of your feature map, to get output with width of 28, you can do upsampling with factor of 2 and upsampling 2 times.

Discriminator:

You will get full credit if you can produce an output of `[batchsize, 1]` vector (image) from the given input `[batchsize, 1, 28, 28]` vector (noise).

Linear Layers that you may use:

- `torch.nn.Conv2d`
- `torch.nn.Linear`

Non-linear Layers:

- `torch.nn.LeakyReLU` with `slope=0.2` between all linear layers.
- `torch.nn.Sigmoid` for the last layer's activation. Can you explain why do we need this in the code comment?

Use Leaky ReLu as the activation function between all layers, except after the last layer use Sigmoid.

You may use `view` to change the vector size: <https://pytorch.org/docs/stable/generated/torch.Tensor.view.html>

As an example, you may use 2 convolution layer and one linear layer in the discriminator, you can also use other setup. Note that instead of using pooling to downsampling, you may also use `stride=2` in convolution to downsample the feature.

```
[19]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
from torch.autograd import Variable
from torchvision.utils import save_image
import matplotlib.pyplot as plt
import matplotlib.animation as animation
from IPython.display import HTML
import numpy as np
from torch.optim.lr_scheduler import StepLR
import torchvision.utils as vutils
from torch.utils.data import DataLoader, TensorDataset
from scipy import linalg
from scipy.stats import entropy
```

```

import tqdm
import cv2
# image input size
image_size=28

# Setting up transforms to resize and normalize
transform=transforms.Compose([
                                transforms.ToTensor(),
                                ])

# batchsize of dataset
batch_size = 100

# Load MNIST Dataset
gan_train_dataset = datasets.MNIST(root='./MNIST/', train=True,
    ↪transform=transform, download=True)
gan_train_loader = torch.utils.data.DataLoader(dataset=gan_train_dataset,
    ↪batch_size=batch_size, shuffle=True)

```

## 0.4 Model Definition (TODO)

```

[20]: class DCGAN_Generator(nn.Module):
        def __init__(self):
            super(DCGAN_Generator,self).__init__()

            #####
            # Please fill in your code here:
            #####

            self.layers = nn.Sequential(
                nn.ConvTranspose2d(100, 128, 4, 1, 0, bias=False),
                nn.BatchNorm2d(128),
                nn.LeakyReLU(0.2),
                nn.ConvTranspose2d(128, 256, 4, 2, 0, bias=False),
                nn.BatchNorm2d(256),
                nn.LeakyReLU(0.2),
                nn.ConvTranspose2d(256, 128, 4, 2, 2, bias=False),
                nn.BatchNorm2d(128),
                nn.LeakyReLU(0.2),
                nn.ConvTranspose2d(128, 64, 2, 2, 2, bias=False),
                nn.BatchNorm2d(64),
                nn.LeakyReLU(0.2),
                nn.ConvTranspose2d(64, 1, 1, 1, 2, bias=False),
                nn.Tanh()
            )

```

```

def forward(self, input):

    #####
    # Please fill in your code here:
    #####
    out = self.layers(input)

    # Explain why Tanh is needed for the last layer
    # The step is to normalize the input from large range to [-1,1] in case
    ↪ of
    # minimal information cannot be captured by the next discriminator. And
    ↪ use
    # Tanh rather than sigmoid can prevent gradients vanishing.

    return out


class DCGAN_Discriminator(nn.Module):
    def __init__(self):
        super(DCGAN_Discriminator, self).__init__()
        #####
        # Please fill in your code here:
        #####
        # Reference from pytorch tutorial
        self.layers = nn.Sequential(
            nn.Conv2d(1, 64, 4, 2, 1, bias=False),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(0.2),
            nn.Conv2d(64, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2),
            nn.Conv2d(128, 512, 4, 2, 1, bias=False),
            nn.BatchNorm2d(512),
            nn.LeakyReLU(0.2),
            nn.Conv2d(512, 1, 4, 2, 1, bias=False),
            nn.Flatten(),
            nn.Linear(1,1),
            nn.Sigmoid()
        )

    def forward(self, input):

        #####
        # Please fill in your code here:
        #####

```

```

        out = self.layers(input)
        # Explain why Sigmoid is needed for the last layer
        # The step is to normalize the input to [0,1] in case of minimal
        ↪ information
        # cannot be make use of decision function with outputs 'fake' and
        ↪ 'Real'.
        # In classification, we use Sigmoid.

        return out

```

```

# Code that check size
g=DCGAN_Generator()
batchsize=2
z=torch.zeros((batchsize, 100, 1, 1))
out = g(z)
print(out.size()) # You should expect size [batchsize, 1, 28, 28]

```

```

d=DCGAN_Discriminator()
x=torch.zeros((batchsize, 1, 28, 28))
out = d(x)
print(out.size()) # You should expect size [batchsize, 1]

```

torch.Size([2, 1, 28, 28])

torch.Size([2, 1])

GAN loss (TODO)

```

[21]: import torch

def loss_discriminator(D, real, G, noise, Valid_label, Fake_label, criterion,
    ↪ optimizerD):

    '''
    1. Forward real images into the discriminator
    2. Compute loss between Valid_label and dicriminator output on real images
    3. Forward noise into the generator to get fake images
    4. Forward fake images to the discriminator
    5. Compute loss between Fake_label and discriminator output on fake images
    ↪ (and remember to detach the gradient from the fake images using detach()!)
    6. sum real loss and fake loss as the loss_D
    7. we also need to output fake images generate by G(noise) for
    ↪ loss_generator computation
    '''
    #####

```

```

    # Please fill in your code here:
    outputs = D(real).view(-1)
    real_loss = criterion(outputs, Valid_label)
    real_loss.backward()
    fake_imgs = G(noise)
    out = D(fake_imgs.detach()).view(-1)
    fake_loss = criterion(out, Fake_label)
    fake_loss.backward()
    loss_D = real_loss + fake_loss
    #####

    return loss_D, fake_imgs

def loss_generator(netD, netG, fake, Valid_label, criterion, optimizerG):
    '''
    1. Forward fake images to the discriminator
    2. Compute loss between valid labels and discriminator output on fake images
    '''

    #####
    # Please fill in your code here:
    outputs = netD(fake).view(-1)
    loss_G = criterion(outputs, Valid_label)
    loss_G.backward()
    #####

    return loss_G

```

```

[22]: import torchvision.utils as vutils
from torch.optim.lr_scheduler import StepLR
import pdb

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Number of channels
nc = 3
# Size of z latent vector (i.e. size of generator input)
nz = 100

netG = DCGAN_Generator().to(device)
netD = DCGAN_Discriminator().to(device)

from torchsummary import summary
print(summary(netG, (100, 1, 1)))
print(summary(netD, (1, 28, 28)))

```

| Layer (type)       | Output Shape      | Param # |
|--------------------|-------------------|---------|
| ConvTranspose2d-1  | [-1, 128, 4, 4]   | 204,800 |
| BatchNorm2d-2      | [-1, 128, 4, 4]   | 256     |
| LeakyReLU-3        | [-1, 128, 4, 4]   | 0       |
| ConvTranspose2d-4  | [-1, 256, 10, 10] | 524,288 |
| BatchNorm2d-5      | [-1, 256, 10, 10] | 512     |
| LeakyReLU-6        | [-1, 256, 10, 10] | 0       |
| ConvTranspose2d-7  | [-1, 128, 18, 18] | 524,288 |
| BatchNorm2d-8      | [-1, 128, 18, 18] | 256     |
| LeakyReLU-9        | [-1, 128, 18, 18] | 0       |
| ConvTranspose2d-10 | [-1, 64, 32, 32]  | 32,768  |
| BatchNorm2d-11     | [-1, 64, 32, 32]  | 128     |
| LeakyReLU-12       | [-1, 64, 32, 32]  | 0       |
| ConvTranspose2d-13 | [-1, 1, 28, 28]   | 64      |
| Tanh-14            | [-1, 1, 28, 28]   | 0       |

Total params: 1,287,360

Trainable params: 1,287,360

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 3.09

Params size (MB): 4.91

Estimated Total Size (MB): 8.01

None

| Layer (type)  | Output Shape     | Param #   |
|---------------|------------------|-----------|
| Conv2d-1      | [-1, 64, 14, 14] | 1,024     |
| BatchNorm2d-2 | [-1, 64, 14, 14] | 128       |
| LeakyReLU-3   | [-1, 64, 14, 14] | 0         |
| Conv2d-4      | [-1, 128, 7, 7]  | 131,072   |
| BatchNorm2d-5 | [-1, 128, 7, 7]  | 256       |
| LeakyReLU-6   | [-1, 128, 7, 7]  | 0         |
| Conv2d-7      | [-1, 512, 3, 3]  | 1,048,576 |
| BatchNorm2d-8 | [-1, 512, 3, 3]  | 1,024     |
| LeakyReLU-9   | [-1, 512, 3, 3]  | 0         |
| Conv2d-10     | [-1, 1, 1, 1]    | 8,192     |
| Flatten-11    | [-1, 1]          | 0         |
| Linear-12     | [-1, 1]          | 2         |
| Sigmoid-13    | [-1, 1]          | 0         |

Total params: 1,190,274

Trainable params: 1,190,274

Non-trainable params: 0

```
-----  
Input size (MB): 0.00  
Forward/backward pass size (MB): 0.54  
Params size (MB): 4.54  
Estimated Total Size (MB): 5.08  
-----
```

None

## TRAINING

```
[23]: import torchvision.utils as vutils  
from torch.optim.lr_scheduler import StepLR  
import pdb  
  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
  
# Number of channels  
nc = 3  
# Size of z latent vector (i.e. size of generator input)  
nz = 100  
  
# Create the generator and discriminator  
netG = DCGAN_Generator().to(device)  
netD = DCGAN_Discriminator().to(device)  
  
# Initialize BCELoss function  
criterion = nn.BCELoss()  
  
# Create latent vector to test the generator performance  
fixed_noise = torch.randn(36, nz, 1, 1, device=device)  
  
# Establish convention for real and fake labels during training  
real_label = 1  
fake_label = 0  
  
learning_rate = 0.0002  
beta1 = 0.5  
  
# Setup Adam optimizers for both G and D  
  
#####  
# Please fill in your code here:  
  
optimizerD = optim.Adam(netD.parameters(), lr=learning_rate, betas=(beta1, 0.  
→999))  
optimizerG = optim.Adam(netG.parameters(), lr=learning_rate, betas=(beta1, 0.  
→999))
```



```
#####

img_list = []
real_img_list = []
G_losses = []
D_losses = []
iters = 0
num_epochs = 10

def load_param(num_eps):
    model_saved = torch.load('/content/gan_{}.pt'.format(num_eps))
    netG.load_state_dict(model_saved['netG'])
    netD.load_state_dict(model_saved['netD'])

# GAN Training Loop
for epoch in range(num_epochs):
    for i, data in enumerate(gan_train_loader, 0):
        real = data[0].to(device)
        b_size = real.size(0)
        noise = torch.randn(b_size, nz, 1, 1, device=device)

        Valid_label = torch.full((b_size,), real_label, dtype=torch.float,
        ↪device=device)
        Fake_label = torch.full((b_size,), fake_label, dtype=torch.float,
        ↪device=device)

        #####
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
        #####

        #####
        # Please fill in your code here:
        netD.zero_grad()
        loss_D, fake_imgs = loss_discriminator(netD, real, netG, noise,
        ↪Valid_label, Fake_label, criterion, optimizerD)
        optimizerD.step()
        #####

        #####
        # (2) Update G network: maximize log(D(G(z)))
        #####

        #####
        # Please fill in your code here:
```

```

netG.zero_grad()
loss_G = loss_generator(netD, netG, fake_imgs, Valid_label, criterion,
↳optimizerG)
optimizerG.step()
#####

# Output training stats
if i % 50 == 0:
    print('%d/%d [%d/%d] \t Loss_D: %.4f \t Loss_G: %.4f \t'
          % (epoch, num_epochs, i, len(gan_train_loader),
             loss_D.item(), loss_G.item()))

# Save Losses for plotting later
G_losses.append(loss_G.item())
D_losses.append(loss_D.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i ==
↳len(gan_train_loader)-1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

    iters += 1

plt.title("Generator and Discriminator Loss During Training")
plt.plot(G_losses, label="G")
plt.plot(D_losses, label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
plt.show()

checkpoint = {'netG': netG.state_dict(),
              'netD': netD.state_dict()}
torch.save(checkpoint, 'gan_{}.pt'.format(num_epochs))

```

```


[0/10] [0/600]   Loss_D: 1.4395   Loss_G: 0.8541
[0/10] [50/600]  Loss_D: 0.0224   Loss_G: 5.5533
[0/10] [100/600] Loss_D: 0.0036   Loss_G: 6.5800
[0/10] [150/600] Loss_D: 0.0019   Loss_G: 7.1339
[0/10] [200/600] Loss_D: 0.0157   Loss_G: 8.5549
[0/10] [250/600] Loss_D: 0.0021   Loss_G: 9.4847
[0/10] [300/600] Loss_D: 0.0037   Loss_G: 9.3039

```

|        |           |                |                |
|--------|-----------|----------------|----------------|
| [0/10] | [350/600] | Loss_D: 0.0007 | Loss_G: 8.7520 |
| [0/10] | [400/600] | Loss_D: 0.0009 | Loss_G: 8.6666 |
| [0/10] | [450/600] | Loss_D: 1.0989 | Loss_G: 5.1204 |
| [0/10] | [500/600] | Loss_D: 0.3220 | Loss_G: 3.9213 |
| [0/10] | [550/600] | Loss_D: 0.6191 | Loss_G: 2.2830 |
| [1/10] | [0/600]   | Loss_D: 0.3472 | Loss_G: 2.8101 |
| [1/10] | [50/600]  | Loss_D: 0.2995 | Loss_G: 2.0540 |
| [1/10] | [100/600] | Loss_D: 0.3558 | Loss_G: 2.5007 |
| [1/10] | [150/600] | Loss_D: 1.0873 | Loss_G: 6.6616 |
| [1/10] | [200/600] | Loss_D: 0.3044 | Loss_G: 3.8561 |
| [1/10] | [250/600] | Loss_D: 0.2065 | Loss_G: 2.8666 |
| [1/10] | [300/600] | Loss_D: 0.7712 | Loss_G: 6.0092 |
| [1/10] | [350/600] | Loss_D: 0.2105 | Loss_G: 2.6980 |
| [1/10] | [400/600] | Loss_D: 0.3247 | Loss_G: 3.1079 |
| [1/10] | [450/600] | Loss_D: 0.4064 | Loss_G: 2.2099 |
| [1/10] | [500/600] | Loss_D: 0.3653 | Loss_G: 3.6658 |
| [1/10] | [550/600] | Loss_D: 0.7750 | Loss_G: 3.5491 |
| [2/10] | [0/600]   | Loss_D: 0.2083 | Loss_G: 2.7046 |
| [2/10] | [50/600]  | Loss_D: 0.4085 | Loss_G: 1.9068 |
| [2/10] | [100/600] | Loss_D: 0.3106 | Loss_G: 2.8062 |
| [2/10] | [150/600] | Loss_D: 0.4981 | Loss_G: 4.3313 |
| [2/10] | [200/600] | Loss_D: 0.4082 | Loss_G: 2.2721 |
| [2/10] | [250/600] | Loss_D: 0.5657 | Loss_G: 1.9689 |
| [2/10] | [300/600] | Loss_D: 0.3133 | Loss_G: 2.9004 |
| [2/10] | [350/600] | Loss_D: 0.4807 | Loss_G: 2.7801 |
| [2/10] | [400/600] | Loss_D: 0.5428 | Loss_G: 1.6733 |
| [2/10] | [500/600] | Loss_D: 0.1836 | Loss_G: 2.8979 |
| [2/10] | [550/600] | Loss_D: 0.6466 | Loss_G: 2.8735 |
| [3/10] | [0/600]   | Loss_D: 0.3426 | Loss_G: 2.5990 |
| [3/10] | [50/600]  | Loss_D: 1.1325 | Loss_G: 5.9731 |
| [3/10] | [100/600] | Loss_D: 0.3000 | Loss_G: 2.0274 |
| [3/10] | [150/600] | Loss_D: 0.6136 | Loss_G: 3.0581 |
| [3/10] | [200/600] | Loss_D: 1.2000 | Loss_G: 4.4493 |
| [3/10] | [250/600] | Loss_D: 0.5336 | Loss_G: 2.2393 |
| [3/10] | [300/600] | Loss_D: 0.2679 | Loss_G: 2.5504 |
| [3/10] | [350/600] | Loss_D: 0.2076 | Loss_G: 2.8943 |
| [3/10] | [400/600] | Loss_D: 0.2621 | Loss_G: 2.7931 |
| [3/10] | [450/600] | Loss_D: 0.5845 | Loss_G: 2.0853 |
| [3/10] | [500/600] | Loss_D: 1.1475 | Loss_G: 3.7529 |
| [3/10] | [550/600] | Loss_D: 0.4137 | Loss_G: 3.5502 |
| [4/10] | [0/600]   | Loss_D: 0.4417 | Loss_G: 1.5140 |
| [4/10] | [50/600]  | Loss_D: 0.3532 | Loss_G: 2.8051 |
| [4/10] | [100/600] | Loss_D: 0.2610 | Loss_G: 2.4036 |
| [4/10] | [150/600] | Loss_D: 0.8951 | Loss_G: 3.4137 |
| [4/10] | [200/600] | Loss_D: 0.3779 | Loss_G: 2.5432 |
| [4/10] | [250/600] | Loss_D: 0.3277 | Loss_G: 2.8436 |
| [4/10] | [300/600] | Loss_D: 0.2879 | Loss_G: 2.5572 |
| [4/10] | [350/600] | Loss_D: 0.2485 | Loss_G: 3.4861 |

|        |           |                |                |
|--------|-----------|----------------|----------------|
| [4/10] | [400/600] | Loss_D: 0.6790 | Loss_G: 1.5138 |
| [4/10] | [450/600] | Loss_D: 0.2840 | Loss_G: 3.3710 |
| [4/10] | [500/600] | Loss_D: 0.4320 | Loss_G: 2.0600 |
| [4/10] | [550/600] | Loss_D: 0.3722 | Loss_G: 1.7186 |
| [5/10] | [0/600]   | Loss_D: 0.3225 | Loss_G: 3.4825 |
| [5/10] | [50/600]  | Loss_D: 0.4497 | Loss_G: 3.5879 |
| [5/10] | [100/600] | Loss_D: 0.2294 | Loss_G: 3.1065 |
| [5/10] | [150/600] | Loss_D: 0.4005 | Loss_G: 3.5037 |
| [5/10] | [200/600] | Loss_D: 0.4449 | Loss_G: 1.8653 |
| [5/10] | [250/600] | Loss_D: 0.5173 | Loss_G: 3.6596 |
| [5/10] | [300/600] | Loss_D: 0.5166 | Loss_G: 1.4815 |
| [5/10] | [350/600] | Loss_D: 0.2562 | Loss_G: 2.3746 |
| [5/10] | [400/600] | Loss_D: 0.2232 | Loss_G: 2.8050 |
| [5/10] | [450/600] | Loss_D: 0.3722 | Loss_G: 3.1250 |
| [5/10] | [500/600] | Loss_D: 0.2260 | Loss_G: 2.9418 |
| [5/10] | [550/600] | Loss_D: 0.2482 | Loss_G: 2.2199 |
| [6/10] | [0/600]   | Loss_D: 0.3549 | Loss_G: 2.3177 |
| [6/10] | [50/600]  | Loss_D: 0.6400 | Loss_G: 3.6243 |
| [6/10] | [100/600] | Loss_D: 0.4446 | Loss_G: 2.6129 |
| [6/10] | [150/600] | Loss_D: 0.2803 | Loss_G: 3.7779 |
| [6/10] | [200/600] | Loss_D: 0.6343 | Loss_G: 2.9019 |
| [6/10] | [250/600] | Loss_D: 0.7729 | Loss_G: 2.2620 |
| [6/10] | [300/600] | Loss_D: 0.2823 | Loss_G: 3.3544 |
| [6/10] | [350/600] | Loss_D: 0.3648 | Loss_G: 2.5363 |
| [6/10] | [400/600] | Loss_D: 0.4429 | Loss_G: 2.9459 |
| [6/10] | [450/600] | Loss_D: 0.3532 | Loss_G: 3.1216 |
| [6/10] | [500/600] | Loss_D: 0.2870 | Loss_G: 2.1404 |
| [6/10] | [550/600] | Loss_D: 0.3004 | Loss_G: 2.5827 |
| [7/10] | [0/600]   | Loss_D: 0.3168 | Loss_G: 3.2149 |
| [7/10] | [50/600]  | Loss_D: 0.3003 | Loss_G: 3.0593 |
| [7/10] | [100/600] | Loss_D: 0.1506 | Loss_G: 5.0919 |
| [7/10] | [150/600] | Loss_D: 0.2250 | Loss_G: 4.0048 |
| [7/10] | [200/600] | Loss_D: 0.2623 | Loss_G: 3.2852 |
| [7/10] | [250/600] | Loss_D: 0.4737 | Loss_G: 3.9448 |
| [7/10] | [300/600] | Loss_D: 0.3490 | Loss_G: 4.1689 |
| [7/10] | [350/600] | Loss_D: 0.9398 | Loss_G: 4.0711 |
| [7/10] | [400/600] | Loss_D: 0.5688 | Loss_G: 1.7997 |
| [7/10] | [450/600] | Loss_D: 0.3657 | Loss_G: 4.6120 |
| [7/10] | [500/600] | Loss_D: 0.1903 | Loss_G: 2.9295 |
| [7/10] | [550/600] | Loss_D: 0.2824 | Loss_G: 3.8851 |
| [8/10] | [0/600]   | Loss_D: 0.2974 | Loss_G: 3.6058 |
| [8/10] | [50/600]  | Loss_D: 0.5625 | Loss_G: 3.7612 |
| [8/10] | [100/600] | Loss_D: 0.2362 | Loss_G: 2.7925 |
| [8/10] | [150/600] | Loss_D: 0.4387 | Loss_G: 3.5062 |
| [8/10] | [200/600] | Loss_D: 0.2416 | Loss_G: 2.5670 |
| [8/10] | [250/600] | Loss_D: 0.2513 | Loss_G: 3.4173 |
| [8/10] | [300/600] | Loss_D: 0.2815 | Loss_G: 3.3367 |
| [8/10] | [350/600] | Loss_D: 0.9523 | Loss_G: 2.4099 |

|        |           |                |                |
|--------|-----------|----------------|----------------|
| [8/10] | [400/600] | Loss_D: 0.8361 | Loss_G: 2.3761 |
| [8/10] | [450/600] | Loss_D: 0.3963 | Loss_G: 2.2499 |
| [8/10] | [500/600] | Loss_D: 0.3237 | Loss_G: 2.5145 |
| [8/10] | [550/600] | Loss_D: 0.2122 | Loss_G: 3.1745 |
| [9/10] | [0/600]   | Loss_D: 0.2427 | Loss_G: 2.8601 |
| [9/10] | [50/600]  | Loss_D: 0.1423 | Loss_G: 4.2237 |
| [9/10] | [100/600] | Loss_D: 0.1019 | Loss_G: 4.2023 |
| [9/10] | [150/600] | Loss_D: 0.3132 | Loss_G: 2.4446 |
| [9/10] | [200/600] | Loss_D: 1.0305 | Loss_G: 3.1943 |
| [9/10] | [250/600] | Loss_D: 0.3721 | Loss_G: 2.8814 |
| [9/10] | [300/600] | Loss_D: 0.5153 | Loss_G: 3.1722 |
| [9/10] | [350/600] | Loss_D: 0.3613 | Loss_G: 3.5833 |
| [9/10] | [400/600] | Loss_D: 0.1698 | Loss_G: 3.7078 |
| [9/10] | [450/600] | Loss_D: 0.4006 | Loss_G: 3.3135 |
| [9/10] | [500/600] | Loss_D: 0.2402 | Loss_G: 3.2695 |
| [9/10] | [550/600] | Loss_D: 0.1962 | Loss_G: 2.6098 |



## 0.5 Qualitative Visualisations

```
[24]: # Test GAN on a random sample and display on 6X6 grid
import matplotlib.pyplot as plt

fig = plt.figure(figsize=(8,8))
```

```
plt.axis("off")
ims = [[plt.imshow(np.transpose(i,(1,2,0)), animated=True)] for i in img_list]
ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000,
    ↳blit=True)

HTML(ani.to_jshtml())
```

[24]: <IPython.core.display.HTML object>

